

Readings

Date	Topic	Required readings	Optional	Due
9/29	<i>Introduction</i>	<ul style="list-style-type: none"> • Allen Institute for AI (2021) • Bonde and Firenze (2013) • Gupta and Heath (2020) 	<ul style="list-style-type: none"> • D. Raji (2019) 	
10/4	<i>Algorithmic decision making</i>	<ul style="list-style-type: none"> • Jiang (2021) • ACLU of Washington (2021) • Kearns and Roth (2019) • Lum and Chowdhury (2021) • stabilityai (n.d.) 		
10/6	<i>What is algorithmic bias?</i>	<ul style="list-style-type: none"> • Baio (2022) • Bembeneck, Nissan, and Obermeyer (2021) • Benjamin (2019) • Chowdhury (2021) • Vox (2021) 	<ul style="list-style-type: none"> • Friedman and Nissenbaum (1996) 	

Date	Topic	Required readings	Optional	Due
10/11	<i>Identifying algorithmic harms</i>	<ul style="list-style-type: none"> • Jeffries, Yin, and Mattu (2020) • Kirchner (2021) • Kirchner (2020) • Varner and Sankin (2020) 		
10/13	<i>Algorithmic fairness</i>	<ul style="list-style-type: none"> • Feathers (2021) • Mitchell et al. (2021) 		
10/18	<i>Ethics of data collection</i>	<ul style="list-style-type: none"> • Smith (2021) • McVean (2019) • Whitby (2020) • Onuoha (2020) 	<ul style="list-style-type: none"> • Ochigame (2020) 	
10/20	<i>Facial recognition</i>	<ul style="list-style-type: none"> • Buolamwini et al. (2020) • Hill and Krolík (2019) • Buolamwini and Gebru (2018) 	<ul style="list-style-type: none"> • Castelvechi (2020) • Noorden (2020) 	
10/25	<i>Criminal Justice</i>	<ul style="list-style-type: none"> • MIT Media Lab (2018) • Angwin et al. (2016) • Hill (2020) • Lum and Isaac (2016) • O'Neill (2019) 		
10/27	<i>Quantifying fairness</i>	<ul style="list-style-type: none"> • Chouldechova (2017) • Mitchell et al. (2021) 	<ul style="list-style-type: none"> • Arvind Narayanan (2018) 	

Date	Topic	Required readings	Optional	Due
11/1	<i>Quantifying fairness</i>	<ul style="list-style-type: none"> • Rudin, Wang, and Coker (2020) • Chouldechova (2020) • Jackson and Mendoza (2020) 		
11/3	<i>Public algorithms</i>	<ul style="list-style-type: none"> • Crawford (2021) • Green (2019) 		
11/8	<i>Private algorithms</i>	<ul style="list-style-type: none"> • Ingold and Soper (2016) • Raghavan and Barocas (2019) 	<ul style="list-style-type: none"> • MIT Technology Review (2021) 	
11/10	<i>Labor</i>	<ul style="list-style-type: none"> • Booth (2020) • Crawford (2021) • Gurley (2021) • The Financial Times (2020) 		
11/15	<i>Welfare</i>	<ul style="list-style-type: none"> • Eubanks (2018) • Henriques-Gomes (2019) • Sudhir and Sunder (2020) 	<ul style="list-style-type: none"> • Gilman (2020) 	
11/17	<i>Healthcare</i>	<ul style="list-style-type: none"> • Chen et al. (2021) • Obermeyer et al. (2019) 		
11/22	<i>Privacy</i>	<ul style="list-style-type: none"> • Mervis (2019) • Wezerek and Riper (2020) • Welford (2018) • Wood et al. (2018) 	<ul style="list-style-type: none"> • 	
11/24		NO CLASS		

Date	Topic	Required readings	Optional	Due
11/29	<i>Paths forward</i>	<ul style="list-style-type: none"> • Green (2020) • Hooker (2021) • I. D. Raji and Buolamwini (2019) • Upchurch (2018) 		
12/1	<i>Paths forward</i>	<ul style="list-style-type: none"> • Mullainathan (2019) • Campbell (2018) • Heilweil (2020) • Nickelsburg (2021) • Richardson (2019) 		
12/6				Final Presentations
12/8				Final Presentations

ACLU of Washington. 2021. “Automated Decision Making Systems Are Making Some of the Most Important Life Decisions For You, but You Might Not Even Know It.” <https://www.aclu-wa.org/story/automated-decision-making-systems-are-making-some-most-important-life-decisions-you-you-might>.

Allen Institute for AI. 2021. “Ask Delphi.” <https://delphi.allenai.org/>.

Angwin, Julia, Jeff Larson, Surya Mattu, and Lauren Kirchner. 2016. “Machine Bias.” <https://www.propublica.org/article/machine-bias-risk-assessments-in-criminal-sentencing?token=jRdnwabwdw5HLiHY-R3nqWS5DOjEM7W->.

Arvind Narayanan. 2018. “Tutorial: 21 Fairness Definitions and Their Politics,” March. <https://www.youtube.com/watch?v=jIXIuYdnyyk>.

Baio, Andy. 2022. “Exploring 12 Million of the 2.3 Billion Images Used to Train Stable Diffusion’s Image Generator.” <https://waxy.org/2022/08/exploring-12-million-of-the-images-used-to-train-stable-diffusions-image-generator/>.

Bembeneck, Emily, Rebecca Nissan, and Ziad Obermeyer. 2021. “To Stop Algorithmic Bias, We First Have to Define It.” <https://www.brookings.edu/research/to-stop-algorithmic-bias-we-first-have-to-define-it/>.

Benjamin, Ruha. 2019. “Default Discrimination.” In *Race After Technology: Abolitionist Tools for the New Jim Code*. Newark, UNITED KINGDOM: Polity Press. <http://ebookcentral.proquest.com/lib/washington/detail.action?docID=5820427>.

Bonde, Sheila, and Paul Firenze. 2013. “A Framework for Making Ethical Decisions.” <https://www.brookings.edu/research/a-framework-for-making-ethical-decisions/>.

- <http://www.brown.edu/academics/science-and-technology-studies/framework-making-ethical-decisions>.
- Booth, Robert. 2020. “Uber Drivers to Launch Legal Bid to Uncover App’s Algorithm.” <http://www.theguardian.com/technology/2020/jul/20/uber-drivers-to-launch-legal-bid-to-uncover-apps-algorithm>.
- Buolamwini, Joy, and Timnit Gebru. 2018. “Conference on Fairness, Accountability and Transparency.” In, 77–91. PMLR. <http://proceedings.mlr.press/v81/buolamwini18a.html>.
- Buolamwini, Joy, Vicente Ordóñez, Jamie Morgenstern, and Erik Learned-Miller. 2020. “Facial Recognition Technologies: A Primer,” May. https://global-uploads.webflow.com/5e027ca188c99e3515b404b7/5ed1002058516c11edc66a14_FRTsPrimerMay2020.pdf.
- Campbell, Alexia Fernández. 2018. “How Tech Employees Are Pushing Silicon Valley to Put Ethics Before Profit.” *Vox*, October. <https://www.vox.com/technology/2018/10/18/17989482/google-amazon-employee-ethics-contracts>.
- Castelvecchi, Davide. 2020. “Is Facial Recognition Too Biased to Be Let Loose?” *Nature* 587 (7834): 347–49. <https://doi.org/10.1038/d41586-020-03186-4>.
- Chen, Irene Y., Emma Pierson, Sherri Rose, Shalmali Joshi, Kadija Ferryman, and Marzyeh Ghassemi. 2021. “Ethical Machine Learning in Healthcare.” *Annual Review of Biomedical Data Science* 4 (1): null. <https://doi.org/10.1146/annurev-biodatasci-092820-114757>.
- Chouldechova, Alexandra. 2017. “Fair Prediction with Disparate Impact: A Study of Bias in Recidivism Prediction Instruments.” *Big Data* 5 (2): 153–63. <https://doi.org/10.1089/big.2016.0047>.
- . 2020. “Transparency and Simplicity in Criminal Risk Assessment.” *Harvard Data Science Review* 2 (1). <https://doi.org/10.1162/99608f92.b9343eec>.
- Chowdhury, Rumman. 2021. “Sharing Learnings about Our Image Cropping Algorithm.” https://blog.twitter.com/engineering/en_us/topics/insights/2021/sharing-learnings-about-our-image-cropping-algorithm.
- Crawford, Kate. 2021. “State.” In, 181–209. *Power, Politics, and the Planetary Costs of Artificial Intelligence*. Yale University Press. <https://doi.org/10.2307/j.ctv1ghv45t.9>.
- Eubanks, Virginia. 2018. “A Child Abuse Prediction Model Fails Poor Families.” *Wired*, January. <https://www.wired.com/story/excerpt-from-automating-inequality/>.
- Feathers, Todd. 2021. “Police Are Telling ShotSpotter to Alter Evidence From Gunshot-Detecting AI.” <https://www.vice.com/en/article/qj8xbq/police-are-telling-shotspotter-to-alter-evidence-from-gunshot-detecting-ai>.
- Friedman, Batya, and Helen Nissenbaum. 1996. “Bias in Computer Systems.” *ACM Transactions on Information Systems* 14 (3): 330347. <https://doi.org/10.1145/230538.230561>.
- Gilman, Michele. 2020. “Poverty Lawgorithms.” <https://datasociety.net/library/poverty-lawgorithms/>.
- Green, Ben. 2019. “The Innovative City: The Relationship Between Technical and Nontechnical Change in City Government.” In *The Smart Enough City: Putting Technology in Its Place to Reclaim Our Urban Future*. <https://direct.mit.edu/books/book/4204/chapter/172388/The-Innovative-City-The-Relationship-between>.
- . 2020. “The False Promise of Risk Assessments: Epistemic Reform and the Limits of

- Fairness.” In *Proceedings of the 2020 Conference on Fairness, Accountability, and Transparency*, 594–606. FAT* ’20. New York, NY, USA: Association for Computing Machinery. <https://doi.org/10.1145/3351095.3372869>.
- Gupta, Abhishek, and Victoria Heath. 2020. “AI Ethics Groups Are Repeating One of Society’s Classic Mistakes.” *MIT Technology Review*, September. <https://www.technologyreview.com/2020/09/14/1008323/ai-ethics-representation-artificial-intelligence-opinion/>.
- Gurley, Lauren Kaori. 2021. “Amazon Drivers Are Instructed to Drive Recklessly to Meet Delivery Quotas.” <https://www.vice.com/en/article/xgxx54/amazon-drivers-are-instructed-to-drive-recklessly-to-meet-delivery-quotas>.
- Heilweil, Rebecca. 2020. “Facebook Is Taking a Hard Look at Racial Bias in Its Algorithms.” *Vox*, July. <https://www.vox.com/recode/2020/7/22/21334051/facebook-news-feed-instagram-algorithm-racial-bias-civil-rights-audit>.
- Henriques-Gomes, Luke. 2019. “The Automated System Leaving Welfare Recipients Cut Off with Nowhere to Turn.” *The Guardian*, October. <http://www.theguardian.com/technology/2019/oct/16/automated-messages-welfare-australia-system>.
- Hill, Kashmir. 2020. “Wrongfully Accused by an Algorithm.” *The New York Times*, June. <https://www.nytimes.com/2020/06/24/technology/facial-recognition-arrest.html>.
- Hill, Kashmir, and Aaron Krolik. 2019. “How Photos of Your Kids Are Powering Surveillance Technology.” *The New York Times*, October. <https://www.nytimes.com/interactive/2019/10/11/technology/flickr-facial-recognition.html>.
- Hooker, Sara. 2021. “Moving Beyond “Algorithmic Bias Is a Data Problem”.” *Patterns* 2 (4). <https://doi.org/10.1016/j.patter.2021.100241>.
- Ingold, David, and Spencer Soper. 2016. “Amazon Doesn’t Consider the Race of Its Customers. Should It?” <http://www.bloomberg.com/graphics/2016-amazon-same-day/>.
- Jackson, Eugenie, and Christina Mendoza. 2020. “Setting the Record Straight: What the COMPAS Core Risk and Need Assessment Is and Is Not.” *Harvard Data Science Review* 2 (1). <https://doi.org/10.1162/99608f92.1b3dadaa>.
- Jeffries, Adrianne, Leon Yin, and Surya Mattu. 2020. “Swinging the Vote?” *The Markup*, February. <https://themarkup.org/google-the-giant/2020/02/26/wheres-my-email>.
- Jiang, Liwei. 2021. “Towards Machine Ethics and Norms.” <https://medium.com/ai2-blog/towards-machine-ethics-and-norms-d64f2bdde6a3>.
- Kearns, Michael, and Aaron Roth. 2019. “Introduction.” In *The Ethical Algorithm: The Science of Socially Aware Algorithm Design*. Oxford, UNITED STATES: Oxford University Press USA - OSO. <http://ebookcentral.proquest.com/lib/washington/detail.action?docID=5905172>.
- Kirchner, Lauren. 2020. “When Zombie Data Costs You a Home.” *The Markup*, October. <https://themarkup.org/locked-out/2020/10/06/zombie-criminal-records-housing-background-checks>.
- . 2021. “Powerful DNA Software Used in Hundreds of Criminal Cases Faces New Scrutiny.” *The Markup*, March. <https://themarkup.org/news/2021/03/09/powerful-dna-software-used-in-hundreds-of-criminal-cases-faces-new-scrutiny>.
- Lum, Kristian, and Rumman Chowdhury. 2021. “What Is an ‘Algorithm’? It Depends Whom

- You Ask.” *MIT Technology Review*, February. <https://www.technologyreview.com/2021/02/26/1020007/what-is-an-algorithm/>.
- Lum, Kristian, and William Isaac. 2016. “To Predict and Serve?” *Significance* 13 (5): 14–19. <https://doi.org/https://doi.org/10.1111/j.1740-9713.2016.00960.x>.
- McVean, Ada. 2019. “40 Years of Human Experimentation in America: The Tuskegee Study.” <https://www.mcgill.ca/oss/article/history/40-years-human-experimentation-america-tuskegee-study>.
- Mervis, Jeffery. 2019. “Can a Set of Equations Keep U.S. Census Data Private?” *ScienceInsider*. <https://www.science.org/content/article/can-set-equations-keep-us-census-data-private>.
- MIT Media Lab. 2018. “Gender Shades,” February. <https://www.youtube.com/watch?v=TWsW1w-BVo>.
- MIT Technology Review. 2021. “Podcast: In Machines We Trust - Hired by an Algorithm,” June. https://www.youtube.com/watch?v=ztcVB_zh_M0.
- Mitchell, Shira, Eric Potash, Solon Barocas, Alexander D’Amour, and Kristian Lum. 2021. “Algorithmic Fairness: Choices, Assumptions, and Definitions.” *Annual Review of Statistics and Its Application* 8 (1): null. <https://doi.org/10.1146/annurev-statistics-042720-125902>.
- Mullainathan, Sendhil. 2019. “Biased Algorithms Are Easier to Fix Than Biased People.” *The New York Times*, December. <https://www.nytimes.com/2019/12/06/business/algorithm-bias-fix.html>.
- Nickelsburg, Monica. 2021. “Washington State Lawmakers Seek to Ban Government from Using Discriminatory AI Tech.” *GeekWire*, February. <https://www.geekwire.com/2021/washington-state-lawmakers-seek-ban-government-using-ai-tech-discriminates/>.
- Noorden, Richard Van. 2020. “The Ethical Questions That Haunt Facial-Recognition Research.” *Nature* 587 (7834): 354–58. <https://doi.org/10.1038/d41586-020-03187-3>.
- O’Neill, James. 2019. “How Facial Recognition Makes You Safer.” *The New York Times*, June. <https://www.nytimes.com/2019/06/09/opinion/facial-recognition-police-new-york-city.html>.
- Obermeyer, Ziad, Brian Powers, Christine Vogeli, and Sendhil Mullainathan. 2019. “Dissecting Racial Bias in an Algorithm Used to Manage the Health of Populations.” *Science* 366 (6464): 447–53. <https://doi.org/10.1126/science.aax2342>.
- Ochigame, Rodrigo. 2020. “The Long History of Algorithmic Fairness.” <https://phenomenalworld.org/analysis/long-history-algorithmic-fairness>.
- Onuoha, Mimi. 2020. “When Proof Is Not Enough.” <https://fivethirtyeight.com/features/w hen-proof-is-not-enough/>.
- Raghavan, Manish, and Solon Barocas. 2019. “Challenges for Mitigating Bias in Algorithmic Hiring.” <https://www.brookings.edu/research/challenges-for-mitigating-bias-in-algorithmic-hiring/>.
- Raji, Deborah. 2019. “That’s Not Fair!” *XRDS: Crossroads, The ACM Magazine for Students* 25 (3): 44–48. <https://doi.org/10.1145/3313127>.
- Raji, Inioluwa Deborah, and Joy Buolamwini. 2019. “Actionable Auditing: Investigating the Impact of Publicly Naming Biased Performance Results of Commercial AI Products.” In,

429435. AIES '19. New York, NY, USA: Association for Computing Machinery. <https://doi.org/10.1145/3306618.3314244>.
- Richardson, Rashida. 2019. "Confronting Black Boxes: A Shadow Report of the New York City Automated Decision System Task Force." <https://ainowinstitute.org/ads-shadowreport-2019.html>.
- Rudin, Cynthia, Caroline Wang, and Beau Coker. 2020. "The Age of Secrecy and Unfairness in Recidivism Prediction." *Harvard Data Science Review* 2 (1). <https://doi.org/10.1162/99608f92.6ed64b30>.
- Smith, Ben. 2021. "How TikTok Reads Your Mind." *The New York Times*, December. <https://www.nytimes.com/2021/12/05/business/media/tiktok-algorithm.html>.
- stabilityai. n.d. "Stable Diffusion - a Hugging Face Space." <https://huggingface.co/spaces/stabilityai/stable-diffusion>.
- Sudhir, K., and Shyam Sunder. 2020. "What Happens When a Billion Identities Are Digitized?" <https://insights.som.yale.edu/insights/what-happens-when-billion-identities-are-digitized>.
- The Financial Times. 2020. "The Uber Game." <https://ig.ft.com/uber-game>.
- Upchurch, Tom. 2018. "To Work for Society, Data Scientists Need a Hippocratic Oath with Teeth." *Wired UK*. <https://www.wired.co.uk/article/data-ai-ethics-hippocratic-oath-cathy-o-neil-weapons-of-math-destruction>.
- Varner, Maddy, and Aaron Sankin. 2020. "Suckers List: How Allstate's Secret Auto Insurance Algorithm Squeezes Big Spenders." *The Markup*, February. <https://themarkup.org/allstates-algorithm/2020/02/25/car-insurance-suckers-list>.
- Vox. 2021. "Are We Automating Racism?" March. <https://www.youtube.com/watch?v=Ok5sKLXqynQ>.
- Wezerek, Gus, and David Van Riper. 2020. "Opinion | Changes to the Census Could Make Small Towns Disappear." *The New York Times*, February. <https://www.nytimes.com/interactive/2020/02/06/opinion/census-algorithm-privacy.html>.
- Whitby, Andrew. 2020. "A Brief History of the Census—and How Covid-19 Could Change It." *Wired*, April. <https://www.wired.com/story/brief-history-census-how-covid-19-could-change-it/>.
- Wolford, Ben. 2018. "What Is GDPR, the EU's New Data Protection Law?" <https://gdpr.eu/what-is-gdpr/>.
- Wood, Alexandra, Micah Altman, Aaron Bembenek, Mark Bun, Marco Gaboardi, James Honaker, Kobbi Nissim, David O'Brien, Thomas Steinke, and Salil Vadhan. 2018. "Differential Privacy: A Primer for a Non-Technical Audience." *SSRN Electronic Journal*. <https://doi.org/10.2139/ssrn.3338027>.